

Dear all,

There is a change in the assignment because of a problem in Assgn 5 we missed carelessly, our fault. Pls note the problem and the solution. Also, given this change is coming on Tuesday, you may not submit a fully working code upto what was told in the intermediate submission this week, as long as most of the code is there (upto the point mentioned), it is ok.

The problem is that though two processes, creating two sockets, will not be able to bind to the same IP-port, the socket id returned by their socket calls can be the same as socket id is on a per-process basis. So the UDP socket id created when a user process makes `m_socket()` call or `m_bind()` call, is not the socket id that init process understands, so S and R cannot use it to send/receive. So to solve it, init process needs to create the socket and bind it when `m_socket()` or `m_bind()` call is made. This needs some additional synchronization. The solution proposed is below. It makes one assumption that only one process will call `socket()/bind()` at one time (that can be removed, but makes the solution bit more complex, so we allow it).

SC has graciously coded it and it works. If you find any boundary problems, let us know. But essentially, you will have to make all calls on the actual UDP socket from init process only.

Note that whenever you access a shared memory, you have to ensure mutual exclusion, that is understood, so not written separately.

Regards,

-AG

-----  
Init process, in addition to what we have given, does this:

1. Create a shared memory structure `SOCK_INFO` with the following fields: `sock_id`, `IP`, `port`, `errno`. All are initialized to 0.
2. Create 2 semaphores `Sem1` and `Sem2`. Both are initialized to 0.
3. The main thread, after creating S and R and doing everything else, does the following:
  - (a) wait on `Sem1`
  - (b) On being signaled, look at `SOCK_INFO`.
  - (c) If all fields are 0, it is a `m_socket` call. Create a UDP socket. Put the socket id returned in the `sock_id` field of `SOCK_INFO`. If error, put -1 in `sock_id` field and `errno` in `errno` field. Signal on `Sem2`.
  - (d) if `sock_id`, `IP`, and `port` are non-zero, it is a `m_bind` call. Make a `bind()` call on the `sock_id` value, with the `IP` and `port` given. If error, reset `sock_id` to -1 in the structure and put `errno` in `errno` field. Signal on `Sem2`.
  - (e) Go back to wait on `Sem1`

`m_socket` call:

Looks at SM to find a free entry as usual. Then signals on `Sem1` and then waits on `Sem2`. On being woken, checks `sock_id` field of `SOCK_INFO`. If -1, return error and set `errno` correctly. If not, put UDP socket id returned in that field in SM table (same as if `m_socket` called `socket()`) and return the index in SM table as usual. In both cases, reset all fields of `SOCK_INFO` to 0.

m\_bind call:

Find the corresponding actual UDP socket id from the SM table. Put the UDP socket ID, IP, and port in SOCK\_INFO table. Signal Sem1. Then wait on Sem2. On being woken, checks sock\_id field of SOCK\_INFO. If -1, return error and set errno correctly. If not return success.

In both cases, reset all fields of SOCK\_INFO to 0.

Rest remains the same.

Assumption. No two processes will make m\_socket or m\_bind calls together (or the signals can go to wrong process. Handling it will require putting process id, checking, or making SOCK\_INFO an array etc. Lets keep things simple).